

OF@TEIN+ Playground: SmartX Box & Inter-connect Provisioning

OF@TEIN+ 1st
Annual Meeting 2018

Muhammad Usman

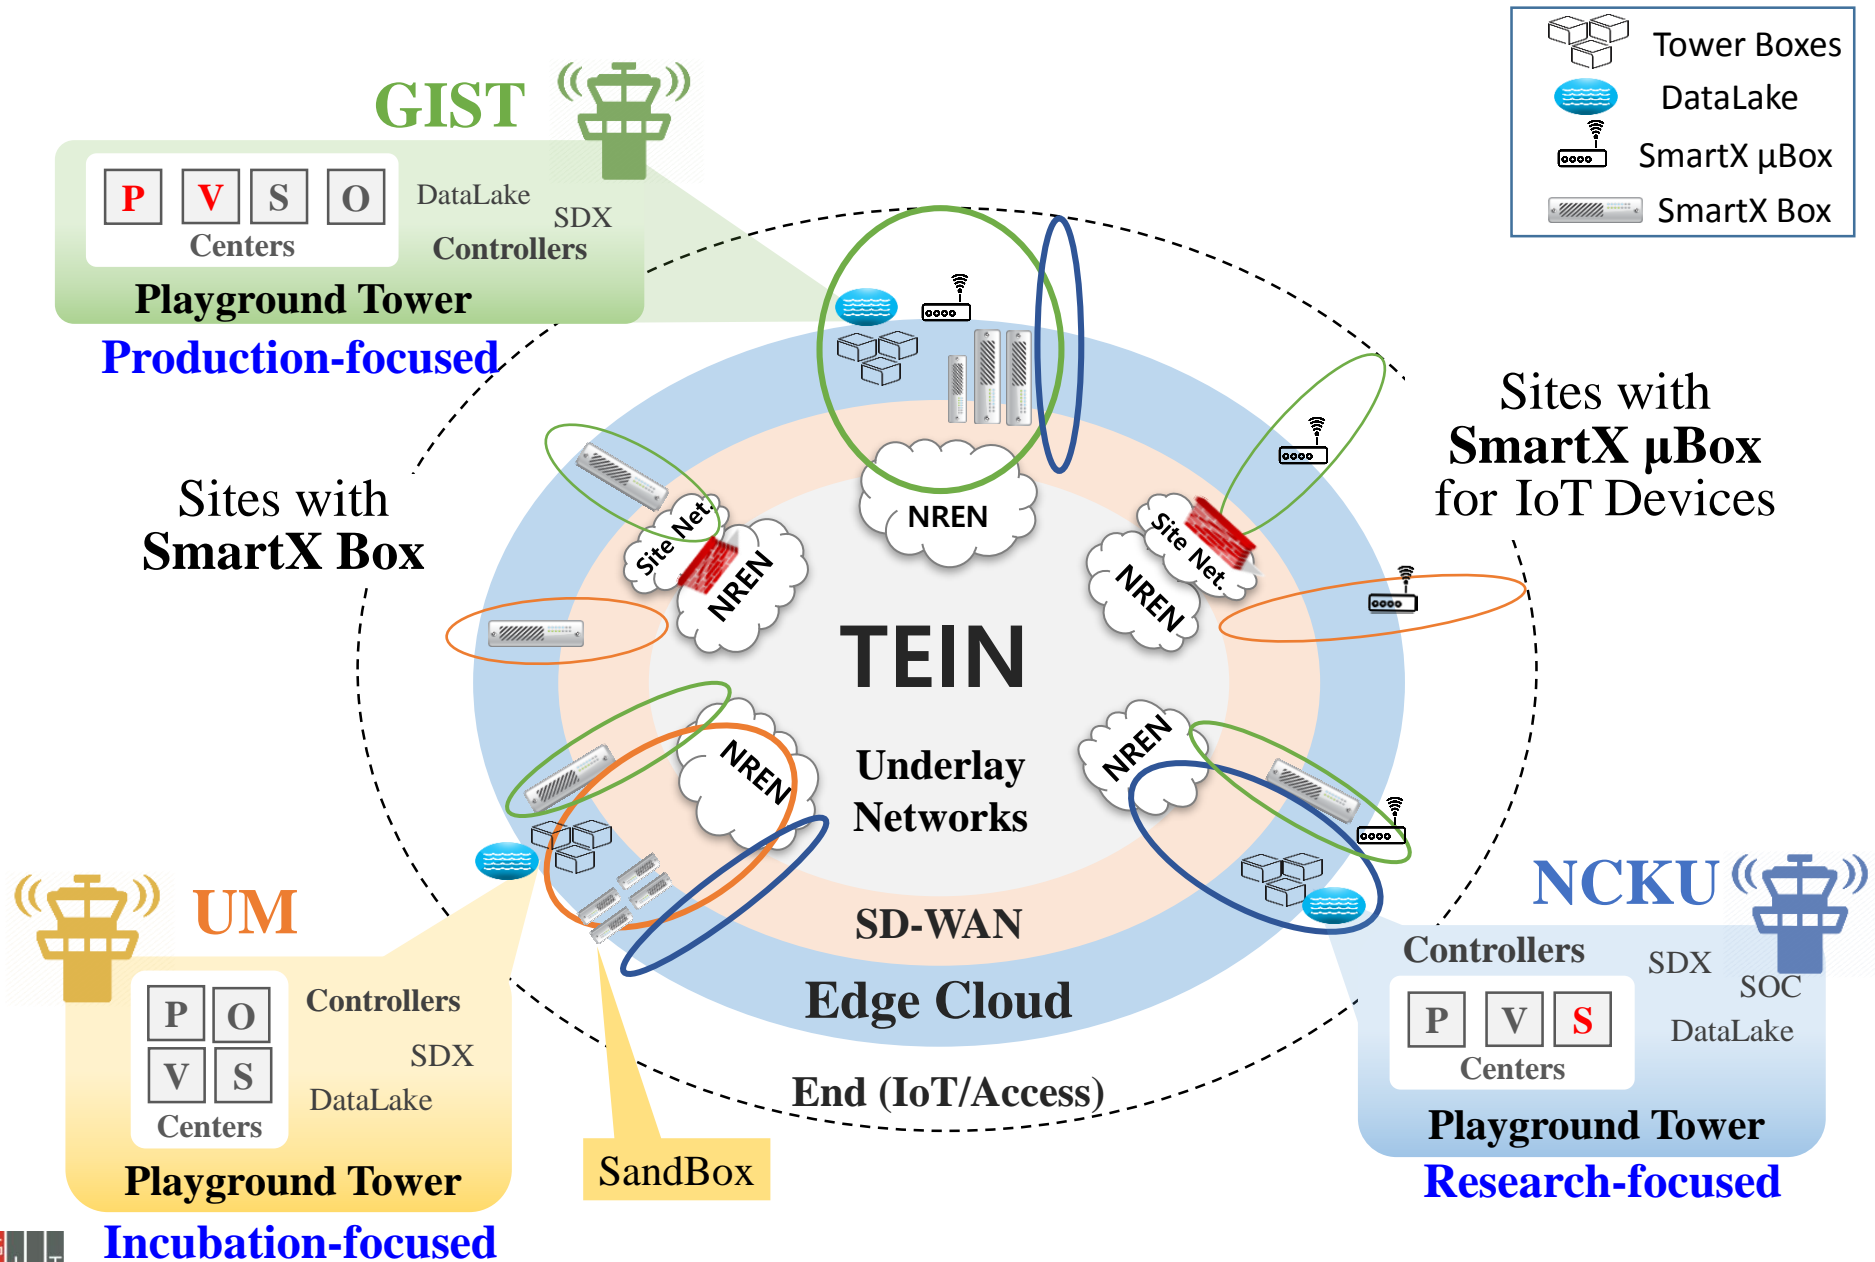
Networked Computing Systems Laboratory (NetCS Lab)
School of Electrical Engineering and Computer Science (EECS)
Gwangju Institute of Science and Technology (GIST)
Gwangju, South Korea

Outline

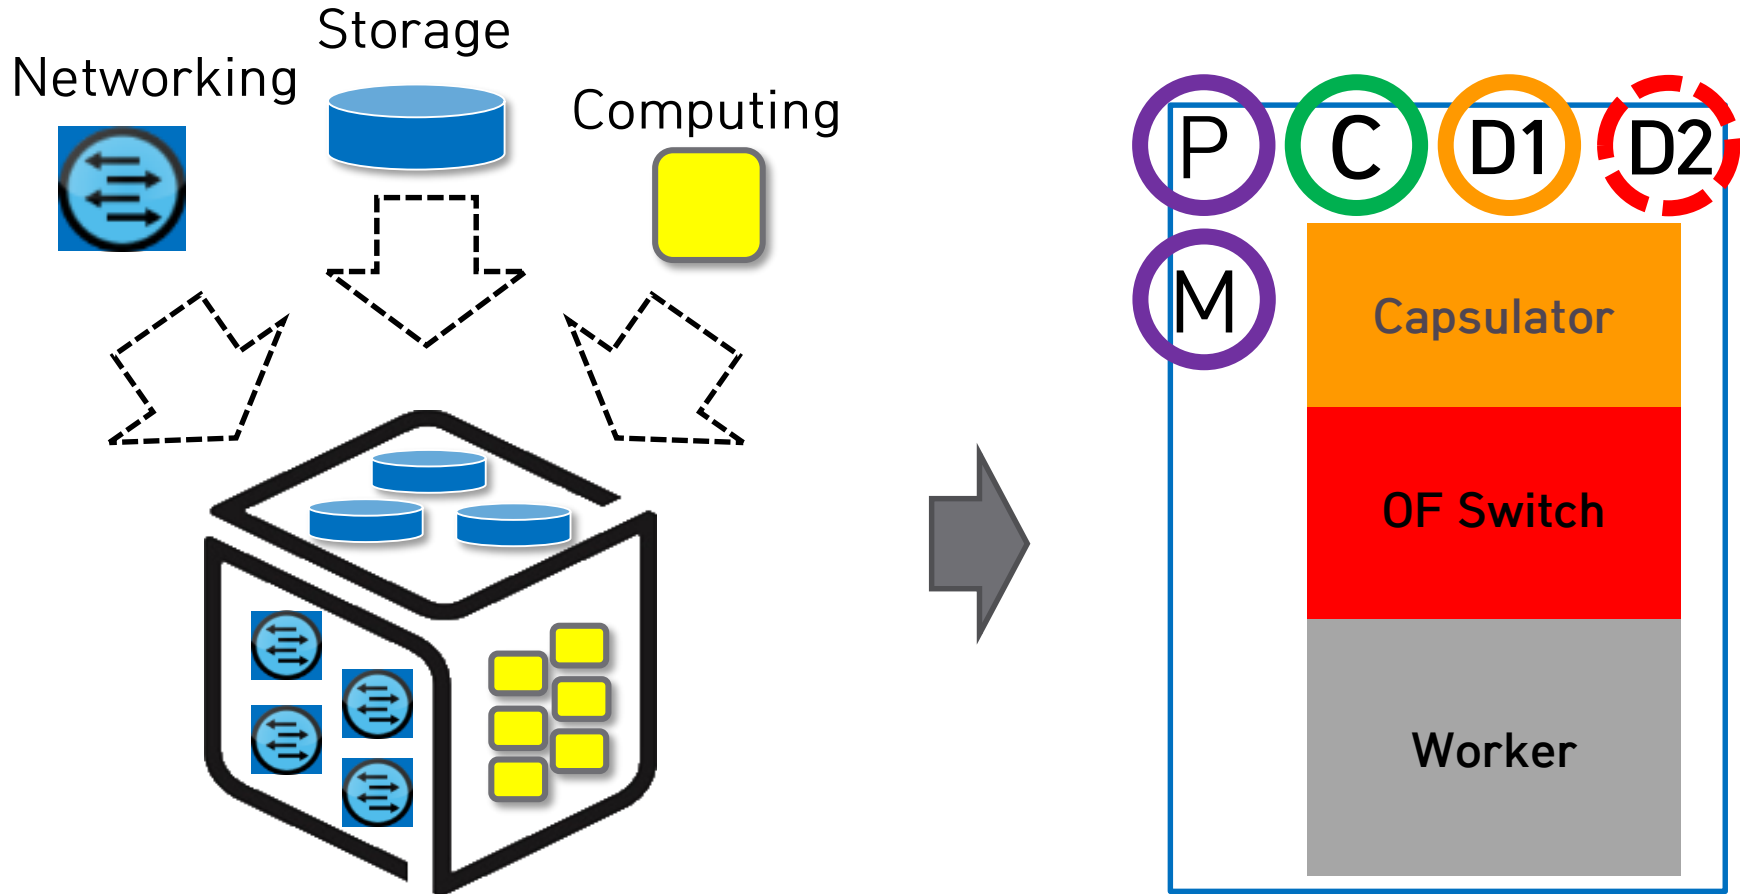
- **OF@TEIN+ Playground: Overview**
 - Draft proposal for OF@TEIN+ Playground
 - SmartX Box concept
- **Step 0: Preparation for Provisioning the SmartX Box**
 - SmartX Box hardware & networking requirements
 - Existing hardware resource types
 - SmartX Box wiring & network connections
- **Step 1: Software Provisioning**
 - SmartX Box key software types
 - SmartX Box 'Triple Star' (***) software design
 - Basic requirements of Triple Star software
 - Software & hardware mapping
 - Triple Star software installation steps
- **Step 2: Inter-connect Provisioning for SmartX Box**
 - SmartX Box software configuration setup
 - SDN inter-connect provisioning
 - Cloud inter-connect provisioning
- **Demonstration**
 - SmartX Box software upgrading for an existing site
 - Video demonstration

Proposal for OF@TEIN+ Playground

2018.01.17



SmartX Box: Hyper-converged Resources



3 Tier Nodes: Capsulator and OF Switch (Networking), Worker (Computing + Storage)
3 Network Planes: Power + Management / Control / Data1 + Data2

Step 0: Preparation for Provisioning SmartX Box

- Hardware & Networking Requirements
- Existing Hardware Types
- Wiring & Network Connections

SmartX Box: Hardware & Networking Requirements

1. SmartX Box Hardware (**Recommended**) Specification

- Processor (CPU) : 8 cores
- Memory (RAM) : 32 GB
- Storage (SSD/HDD) : 512 GB
- Network interface : 5 Ethernet ports (Power + Management + Control + Data1 + Data2 [optional])

2. L2 Switch with dedicated VLANs and subnet IP addresses

- VLAN/Subnet 1 : /28 block for Power + Management + Control
- VLAN/Subnet 2 : /30 block for Data1
- VLAN/Subnet 3 : /29 block for Data2 [optional]

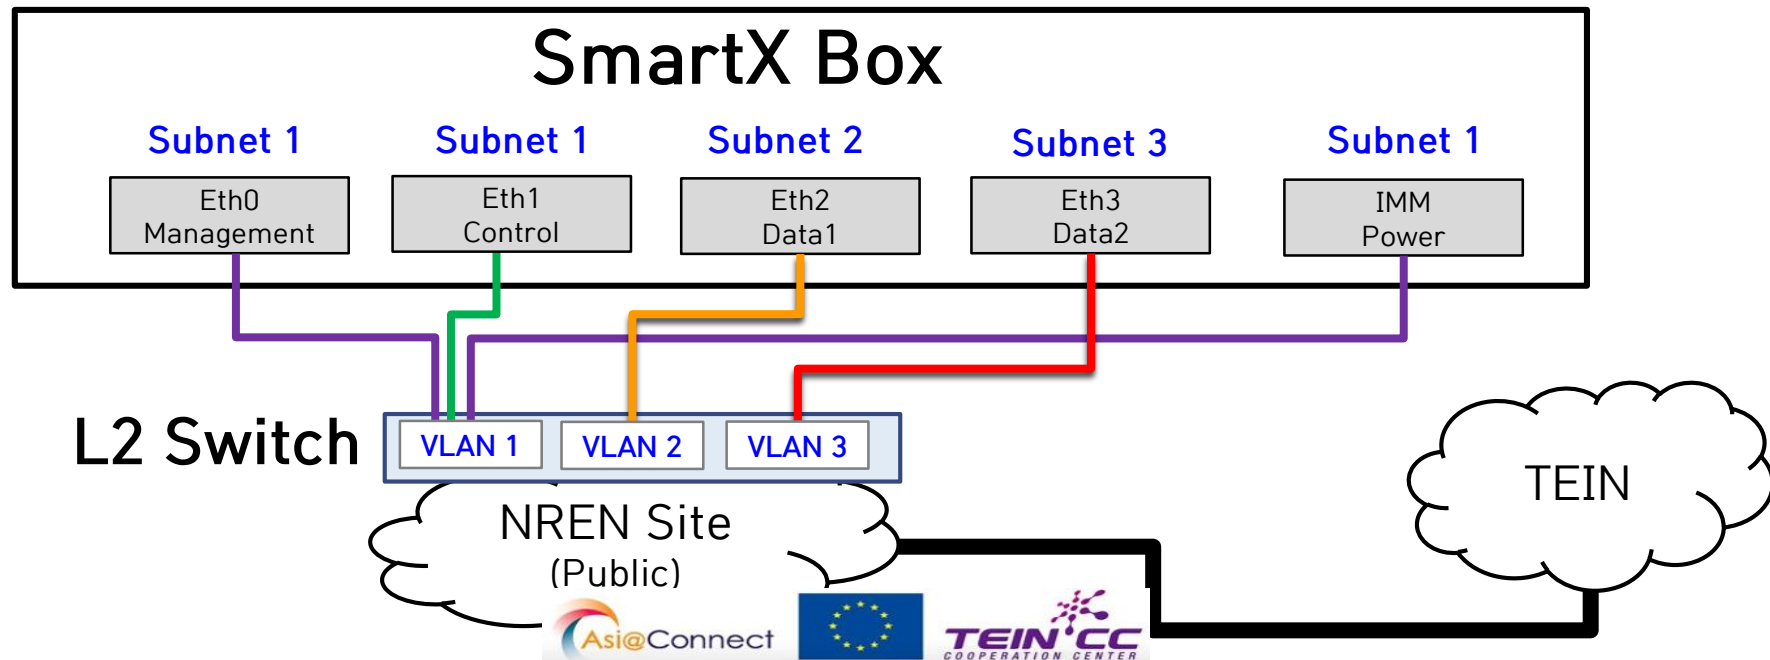
SmartX Box: Hardware Resource Types

Model	Type B (IBM System x3650 M4)	Type O (SuperServer E300-8D)
Size	17.5 in X 29.4in X 3.4in	10in X 8.9in X 8.9in
Processor	Intel® Xeon® processor E5-2630, 12 Cores, 12 Threads, 2.30GHz	Intel® Xeon® processor D-1518, 4 Cores, 8 Threads, 2.2GHz
Memory	24 x DDR3 DIMM sockets (32GB)	4x DDR4 DIMM sockets (32GB)
Graphics	Matrox G200eR2	Aspeed® AST2400
I/O	8 x USB 3.0 ports 2 x VGA port IBM ServerRAID, IBM SATA Device	2 USB 3.0 ports 1 VGA port 4 x SATA3 ports
Network Interfaces	4 x 1GbE LAN ports 2 x 10GbE Embedded Adapter (Optional) 1 x Integrated Management Module II (IMM2)	2 10G SFP+ LAN ports 6 1GbE LAN ports 1 Dedicated IPMI LAN port
Storage	600GB (HDD)	256GB (SSD)



SmartX Box: Wiring & Network Connections

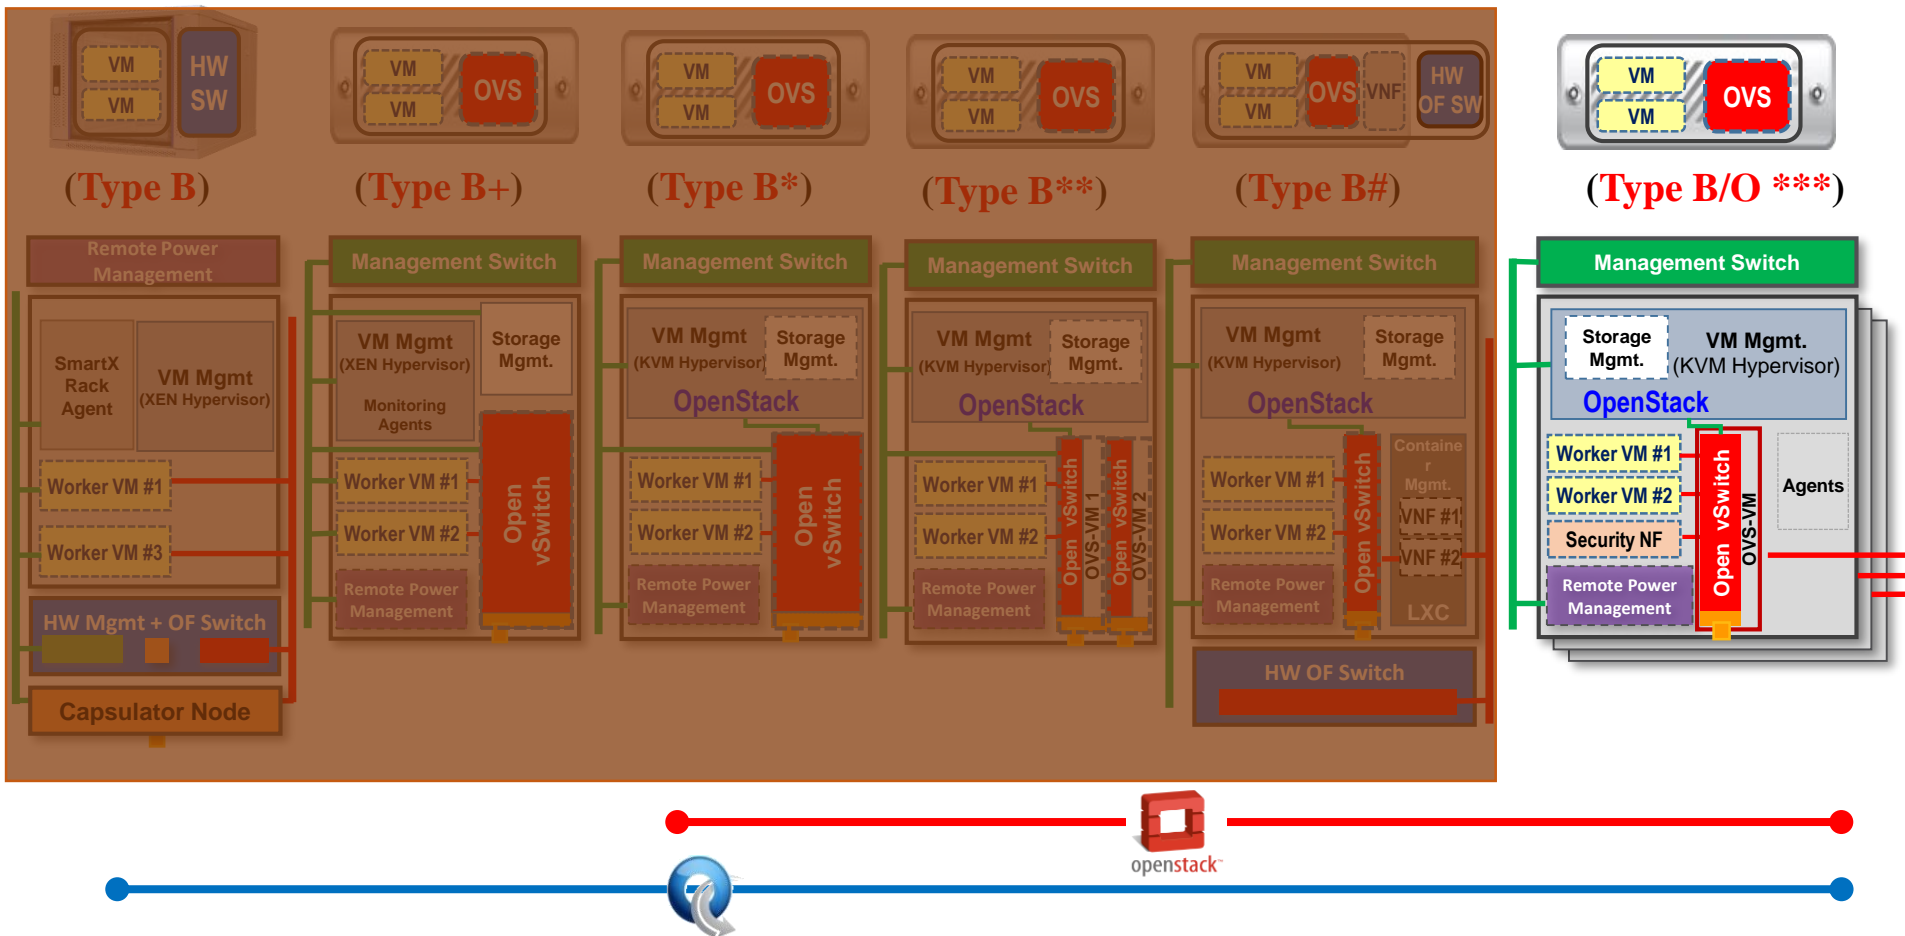
(Example: Type B with IBM M4 Server)



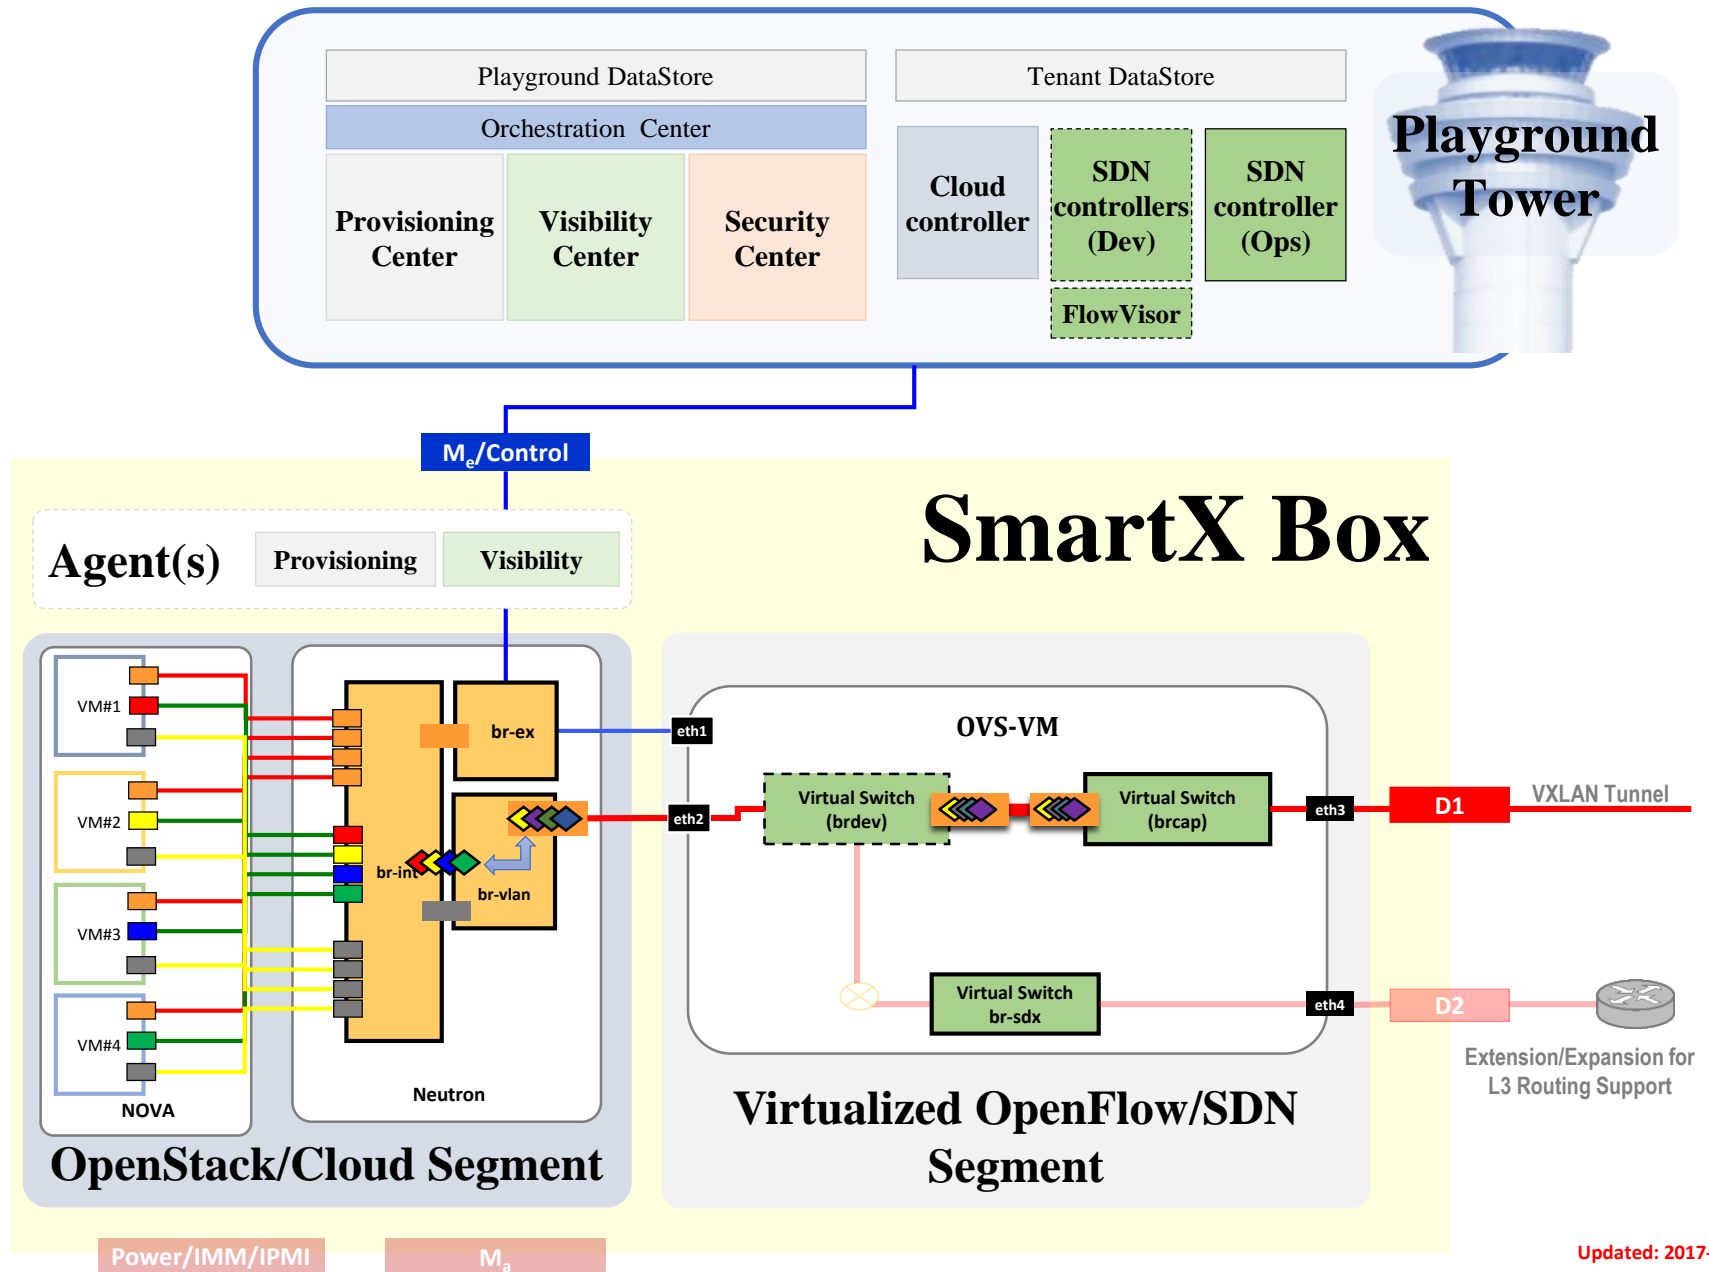
Step 1: Software Provisioning

- Overview
- Software provisioning for Controller's
- Cloud software provisioning for SmartX Box
- SDN software provisioning for SmartX Box

SmartX Box: Key Software Types Timeline



SmartX Box: Triple Star (***) Software Design (v2)



SmartX Box: Basic Requirements of TripleStar (***) Software

Software Component	Version
• Ubuntu version	16.04.03 (LTS)
• Kernel version	4.10.0-28-generic
• VM Management	KVM by OpenStack (Stable Ocata)
• OVS version	2.6.1 (Ubuntu Repository)

SmartX Box: Software & Hardware Mapping

SmartX Box Hardware


Type B



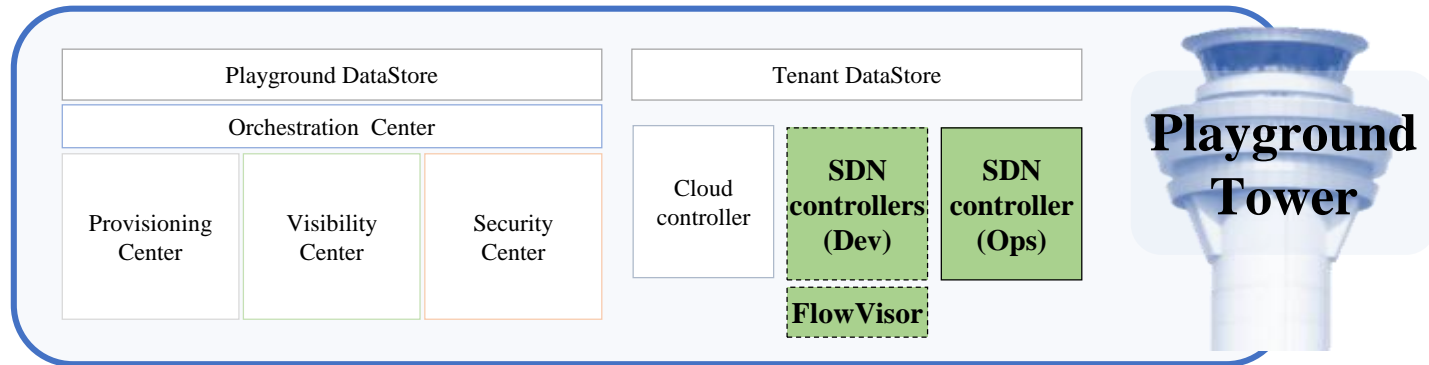
Type O



SmartX Box Software



SDN Controllers Provisioning



Software	Version	Features	Function
FlowVisor		Flow Slicer	Slice or Divide the Flow Table in the Switches
OpenDaylight	Stable Hydrogen	Controller (Dev)	Centralized Policy/Rules Management
ONOS	Stable Loon	Controller (Ops)	Centralized Policy/Rules Management

Flowvisor Provisioning

<https://github.com/OPENNETWORKINGLAB/flowvisor/wiki>

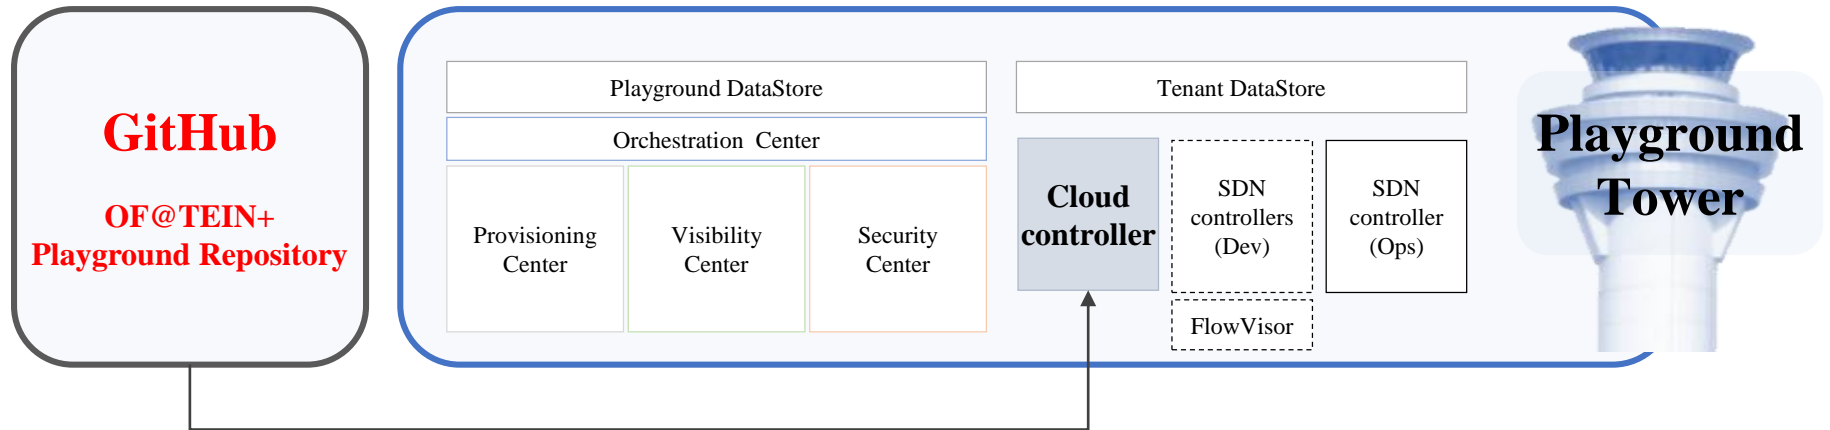
ONOS Controller (Ops) Provisioning

```
$wget http://repo1.maven.org/maven2/org/onosproject/onos-releases/onos-1.11.1/onos-1.11.1.tar.gz
$ tar -xvzf onos-1.11.1.tar.gz
$ cd onos-1.11.1/bin
$ ./onos-service start
```

ODL Controller (Dev) Provisioning

<https://wiki.opendaylight.org/view/Release>

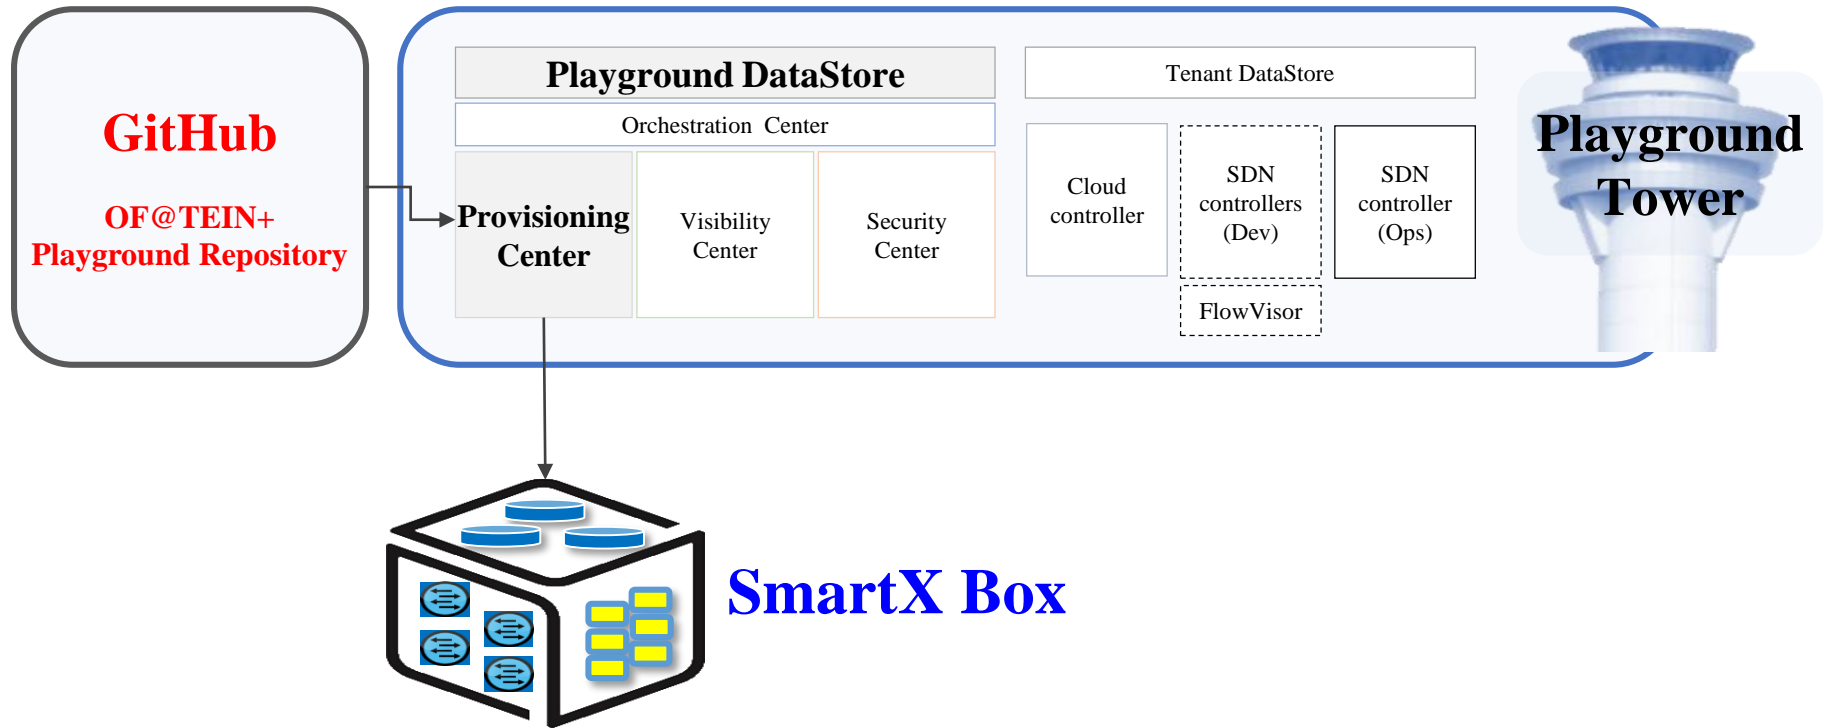
Cloud Controllers Provisioning



Software	Version	Features	Function
OpenStack	Stable Ocata	Keystone	Centralized Account Management
		Horizon	Centralized Dashboard (Web UI)

```
$ git clone https://github.com/OFTEIN-NET/Provisioning-SmartX-Box-TripleStar
$ cd Provisioning-SmartX-Box-TripleStar
$ sudo chmod +x install_controller.sh
$ ./install_controller.sh
```


SmartX Box Provisioning



Software	Version	Features	Function
OpenStack	Stable Ocata	Nova	Instances (VMs) management
		Neutron	Instances Networking Configuration
		Glance	Image Storage
Open vSwitch	2.6.1 (Packages)	OVS Switch	Configuring Operator/Developers SDN Switches
		OVSDB	Centralized management/control SDN switches

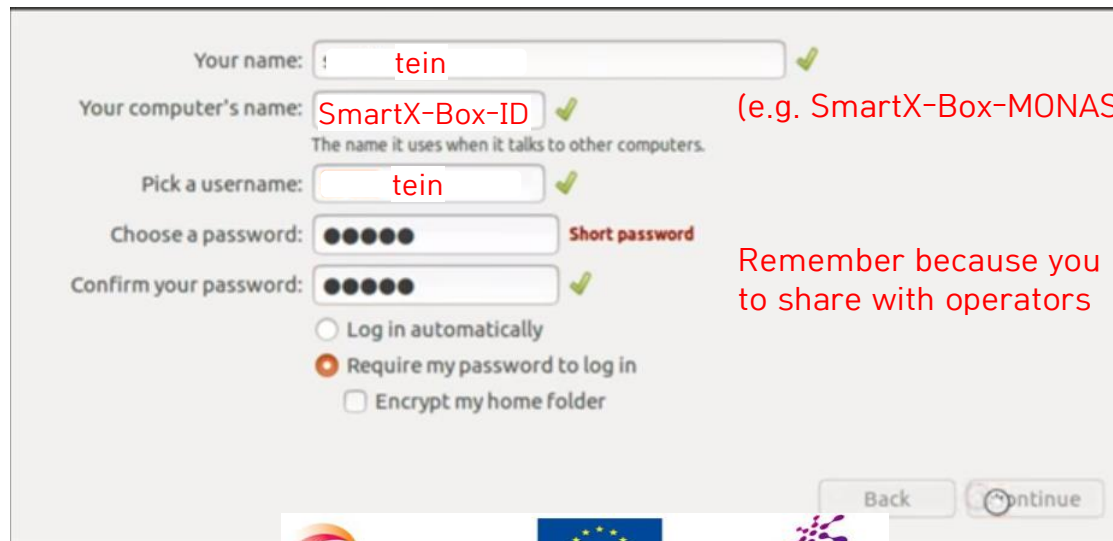
SmartX Box Provisioning: Triple Star (***)

Software Installation Steps

1. Install Ubuntu OS (16.04.3)
2. Check the connectivity to central sites (GIST, MYREN, and NCKU)
3. Submit network-related information to operator
4. Download 'Otriple Star' installation software
5. Execute OpenStack cloud installation script
6. Execute OVS-VM creation script for deploying SDN switches
7. Execute vSwitches creation and configuration script for integrating Cloud and SDN segments
8. Inform and request for an OpenStack account to operator for performing any experiments

(1) Install Ubuntu OS (16.04.3)

- In case you need assistance in installing Ubuntu 16.04.3. Just follow the given link as a reference and modify accordingly
 - https://www.youtube.com/watch?v=c_Ja2PRt0cc
- During Ubuntu installation specify following details



Your name: tein ✓

Your computer's name: SmartX-Box-ID ✓
The name it uses when it talks to other computers.

Pick a username: tein ✓

Choose a password: ●●●● Short password

Confirm your password: ●●●● ✓

☐ Log in automatically

☒ Require my password to log in

☐ Encrypt my home folder

Back Continue

(e.g. SmartX-Box-MONASH for Monash)

Remember because you need to share with operators

(2) Box Connectivity Verification

- Required before Starting installation (if failed, installation will be failed)
- Connectivity to these entities are required:
 - Cloud Controller for OpenStack Centralized Management (Keystone and Horizon) : [103.22.221.74](#)
 - OpenFlow-SDN Network Slicer (FlowVisor) : [103.22.221.52](#)
 - SmartX Configuration and Access Center : [103.22.221.53](#)
 - One of the Testbed HUB nodes (MYREN, NCKU or GIST) : [103.26.47.229](#), [140.116.158.233](#) or [61.252.52.11](#)
- If you are behind a Firewall, please add rules as shown in the next page

(2) Server Connectivity Verification

- Firewall Rules for OpenStack Cloud

OpenStack service	Default ports	Port type
Block Storage (cinder)	8776	publicurl and adminurl
Compute (nova) endpoints	8774	publicurl and adminurl
Compute API (nova-api)	8773, 8775	
Compute ports for access to virtual machine consoles	5900-5999	
Compute VNC proxy for browsers (openstack-nova-novncproxy)	6080	
Compute VNC proxy for traditional VNC clients (openstack-nova-xvncproxy)	6081	
Proxy port for HTML5 console used by Compute service	6082	
Identity service (keystone) administrative endpoint	35357	adminurl
Identity service public endpoint	5000	publicurl
Image Service (glance) API	9292	publicurl and adminurl
Image Service registry	9191	
Networking (neutron)	9696	publicurl and adminurl
Object Storage (swift)	6000, 6001, 6002	
Orchestration (heat) endpoint	8004	publicurl and adminurl
Orchestration AWS CloudFormation-compatible API (openstack-heat-api-cfn)	8000	
Orchestration AWS CloudWatch-compatible API (openstack-heat-api-cloudwatch)	8003	
Telemetry (ceilometer)	8777	publicurl and adminurl

- Compute (nova) placement 8778
- Firewall Rules for OpenFlow SDN
 - OpenvSwitch : VXLAN (UDP/4789), OVSDb (TCP/4455)
 - OpenFlow Controller : OF-Control (TCP/6633), REST API (TCP/8080)

(3) Submit Network Information

- Required for modifying the installation script for the site
- Submit these information to Operator via Email: ops@oftein.net
 - Site Name (Comsats, Monash, ...)
 - Box IP Address (Management IP)
 - Free Public IP Address (for VMs Pool and one IP for OVS instances)
 - Gateway IP Address for Management Subnet
 - Data path subnet details (IP address, netmask, gateway)

**Wait for the confirmation from
OF@TEIN+ Operator(s) before
proceeding to the next steps**

(4) Download 'Triple Star' Software

- Download the Otriple Star software (e.g. when installing at COMSATS University)
 - # scp -r netcs@103.22.221.74:/home/netcs/openstack/Box-Installations/OTripleStar-Install-COMSATS .
- Move inside the downloaded directory
 - # cd [OTripleStar-Install-COMSATS](#)
- Change permissions on the scripts
 - # sudo chmod +x *.sh

(5) Installing OpenStack Cloud Software

- Script Name: `install_smartx_box.sh`
- Execute Installation Script with this command in the Linux CLI
`./install_smartx_box.sh`

(6) Creating OVS-VM for SDN Switches Configurations

- Script Name: `create_ovs_vm.sh`
- Execute Installation Script with this command in the CLI
`./create_ovs_vm.sh`

(7) Creating SDN Switches and Integrating with OpenStack

- Script Name: `create_ovs_bridges.sh`
- Execute SDN switches configuration Script
`./create_ovs_bridges.sh`

(8) Inform + Request Account for Experiment

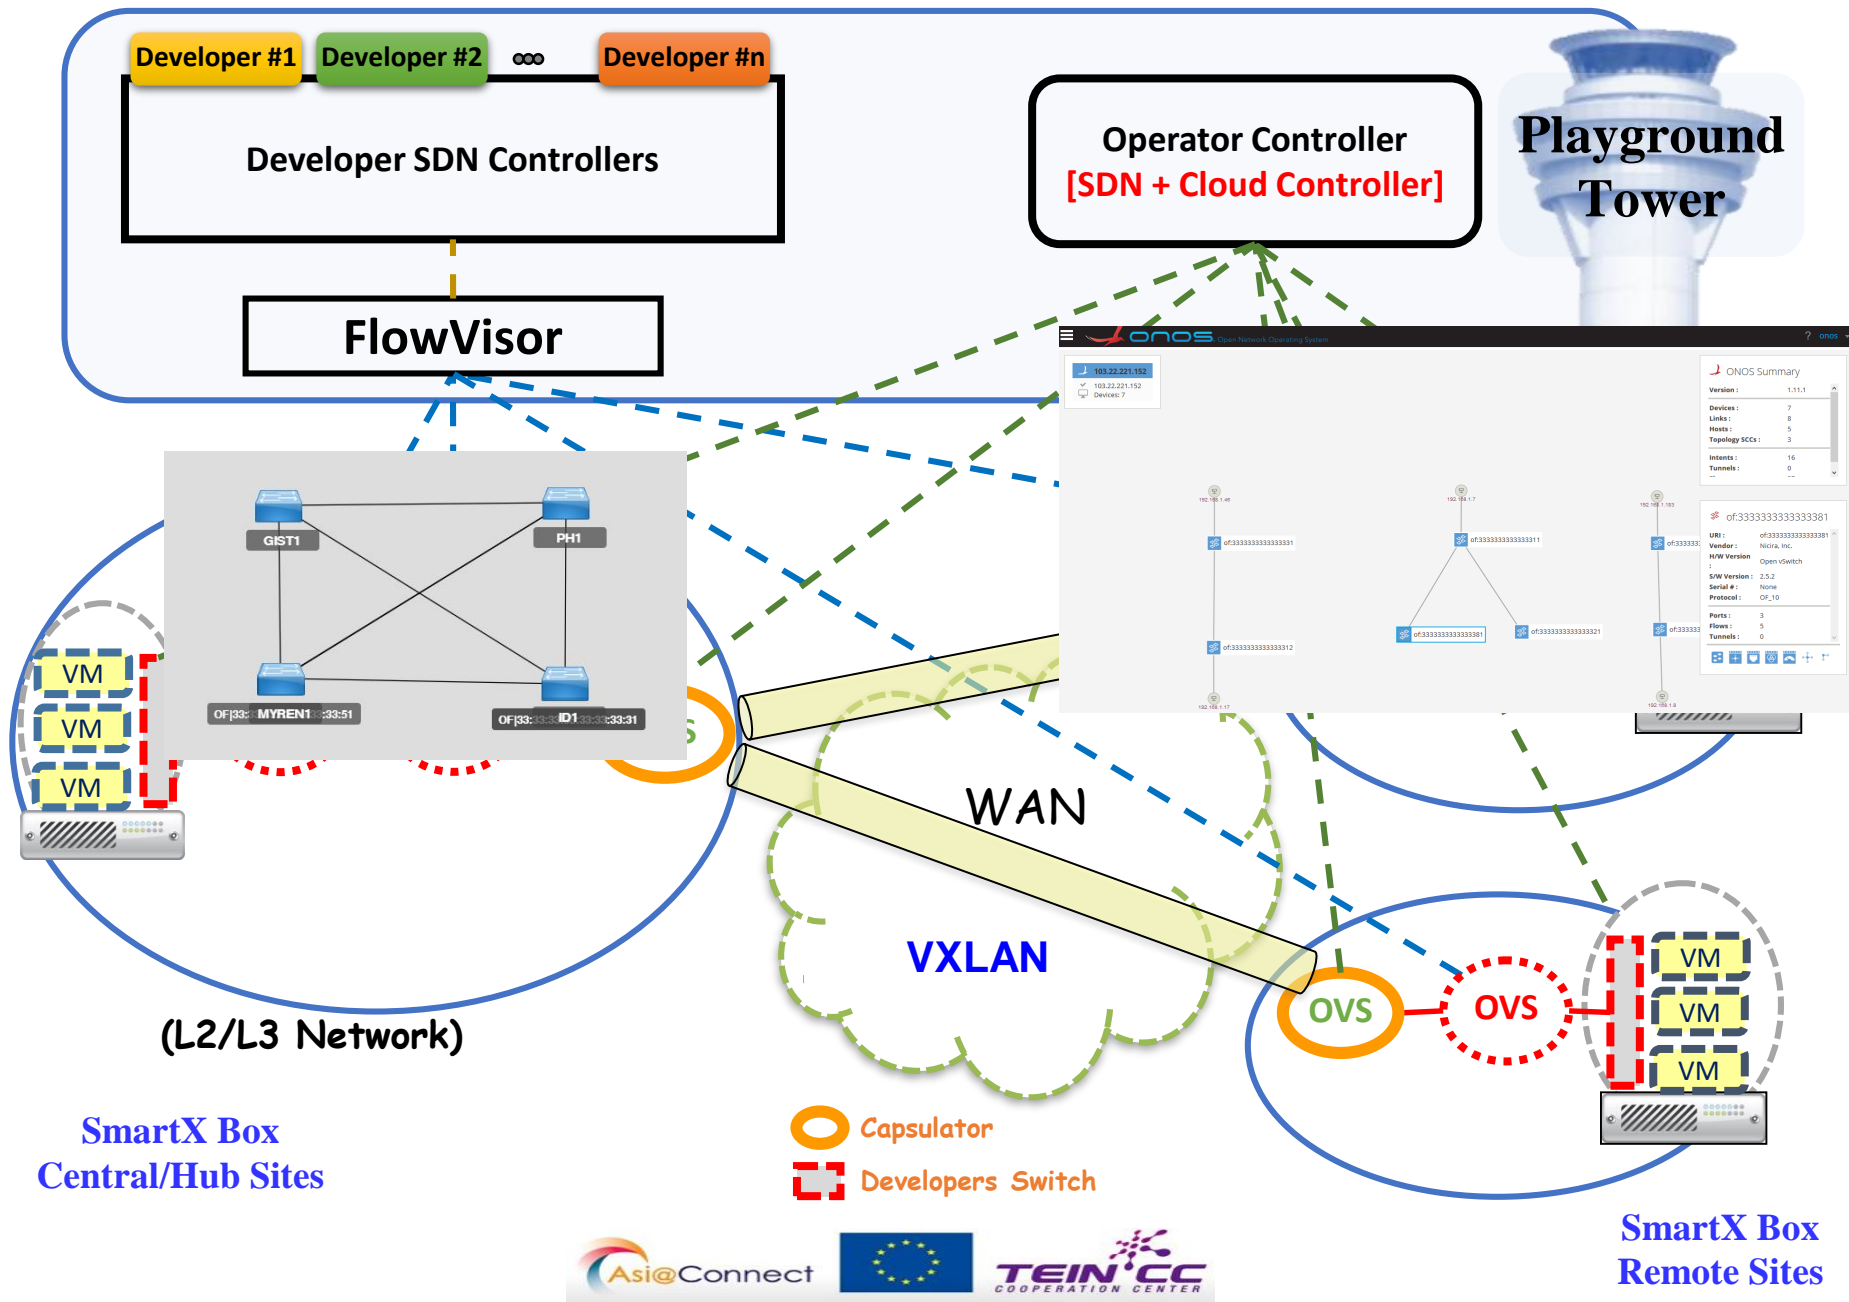
- Send the account request to operator
Email : ops@oftein.net
- Please check detailed information for testing VM-to-VM communication using SDN-Cloud Testbed

<https://www.youtube.com/watch?v=gdxpKvy-75E>

Step 2: Inter-connect Provisioning

- Overview
- Software provisioning for Controller's
- Cloud software provisioning for SmartX Box
- SDN software provisioning for SmartX Box


Multi-islands SDN with Network Slicing



Access Operator Controller & Add ONOS Intents

103.22.221.152:8181

Started



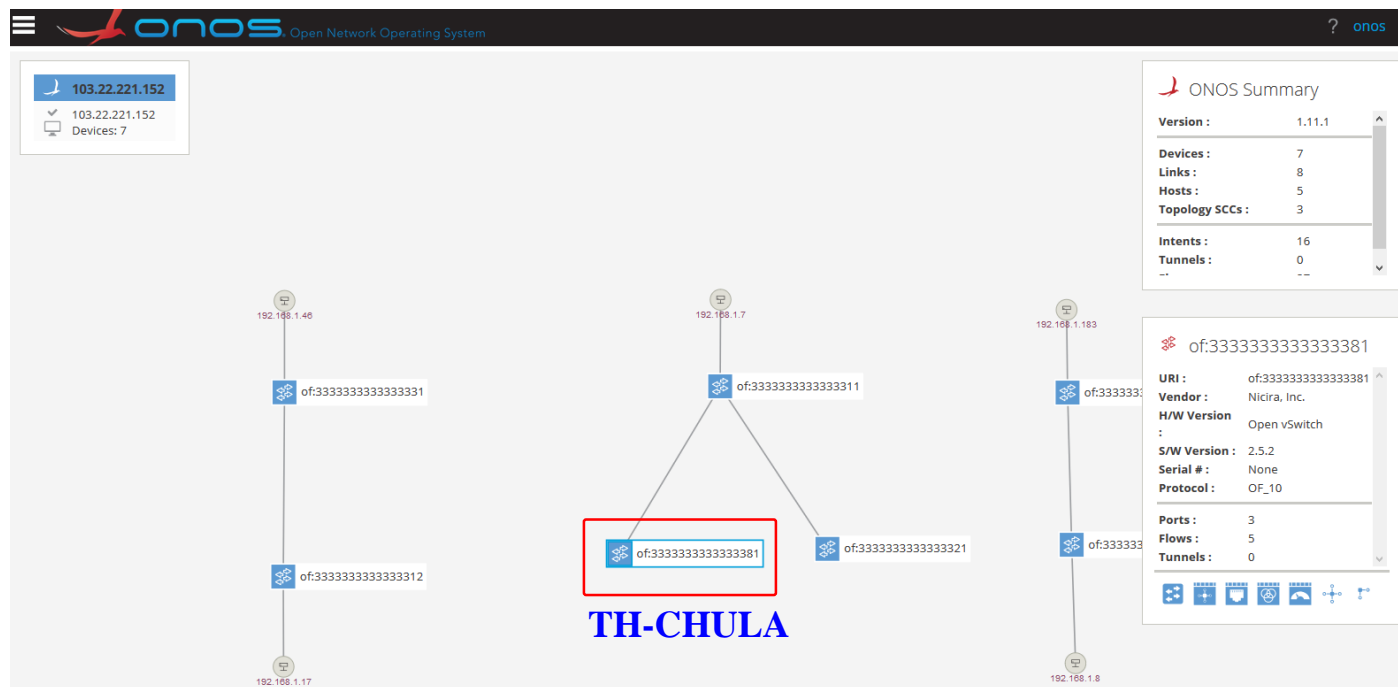
Open Network Operating System

Enter User details

User:

Password:

Login



of:3333333333333381

URI : of:3333333333333381

Vendor : Nicira, Inc.

Type : SWITCH

H/W Version : Open vSwitch

Master ID : 103.22.221.152

S/W Version : 2.5.2

Chassis ID :

Protocol : OF_10

Serial # : None

**Check ports for
adding intents**

Ports

Enabled	ID	Speed	Type	Egress Links	Name
false	Local	0	Copper		brcap
true	1	0	Copper		C_KR_GIST1
true	2	0	Copper	of:3333333333333311/4	ovs_vxlan_GIST1

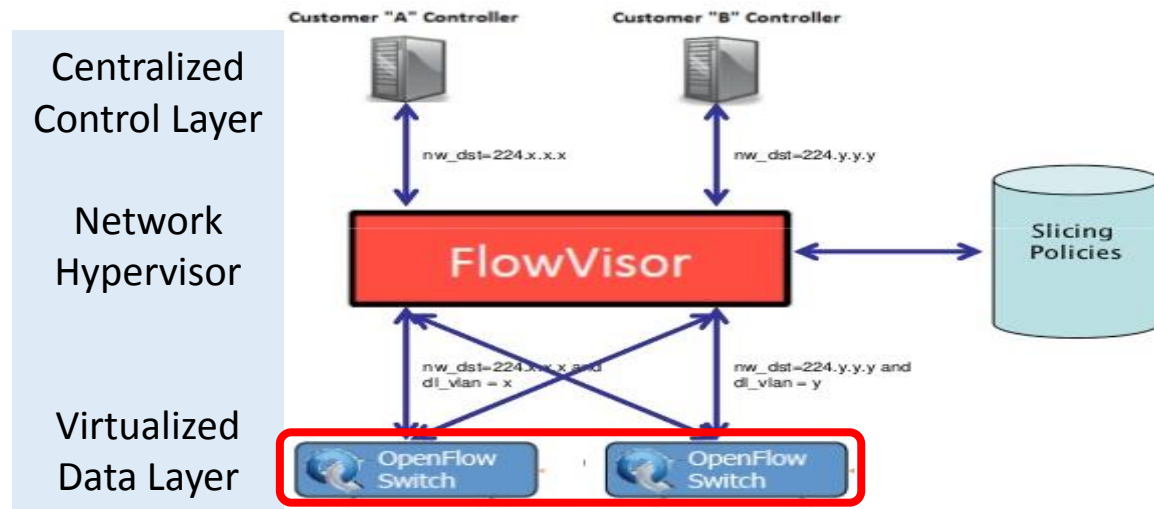
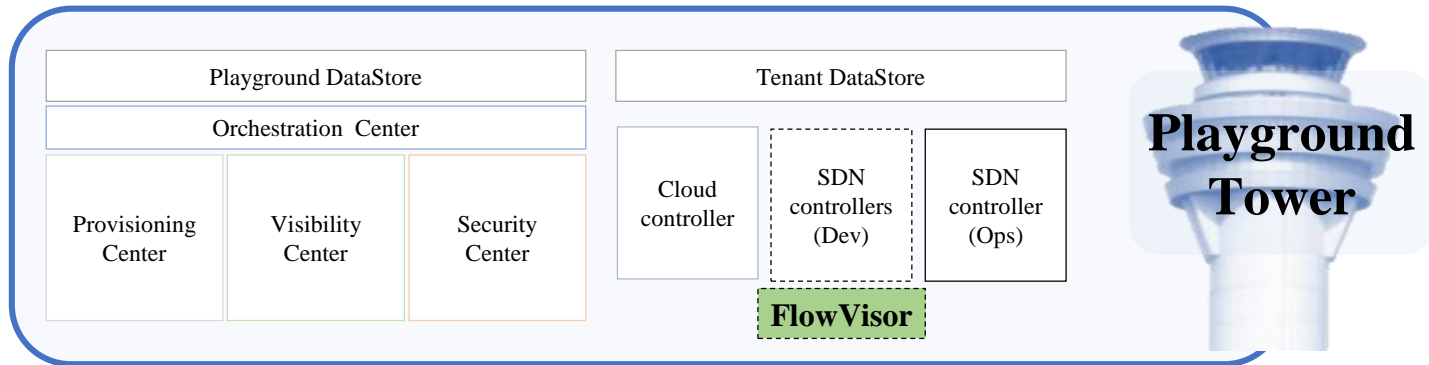
#KR-GIST1

add-point-intent of:3333333333333311/1 of:3333333333333311/2
 add-point-intent of:3333333333333311/2 of:3333333333333311/1
 add-point-intent of:3333333333333311/4 of:3333333333333311/3
 add-point-intent of:3333333333333311/3 of:3333333333333311/4

#TH-CHULA

add-point-intent of:3333333333333381/2 of:3333333333333381/1
 add-point-intent of:3333333333333381/1 of:3333333333333381/2

Virtual Networking VLAN-based network Slicing (FlowVisor + OpenvSwitch)



Adding new slice & flowpace

```
fvctl-json --passwd-file=passwd list-datapaths
```

```
fvctl-json --passwd-file=passwd list-slices
```

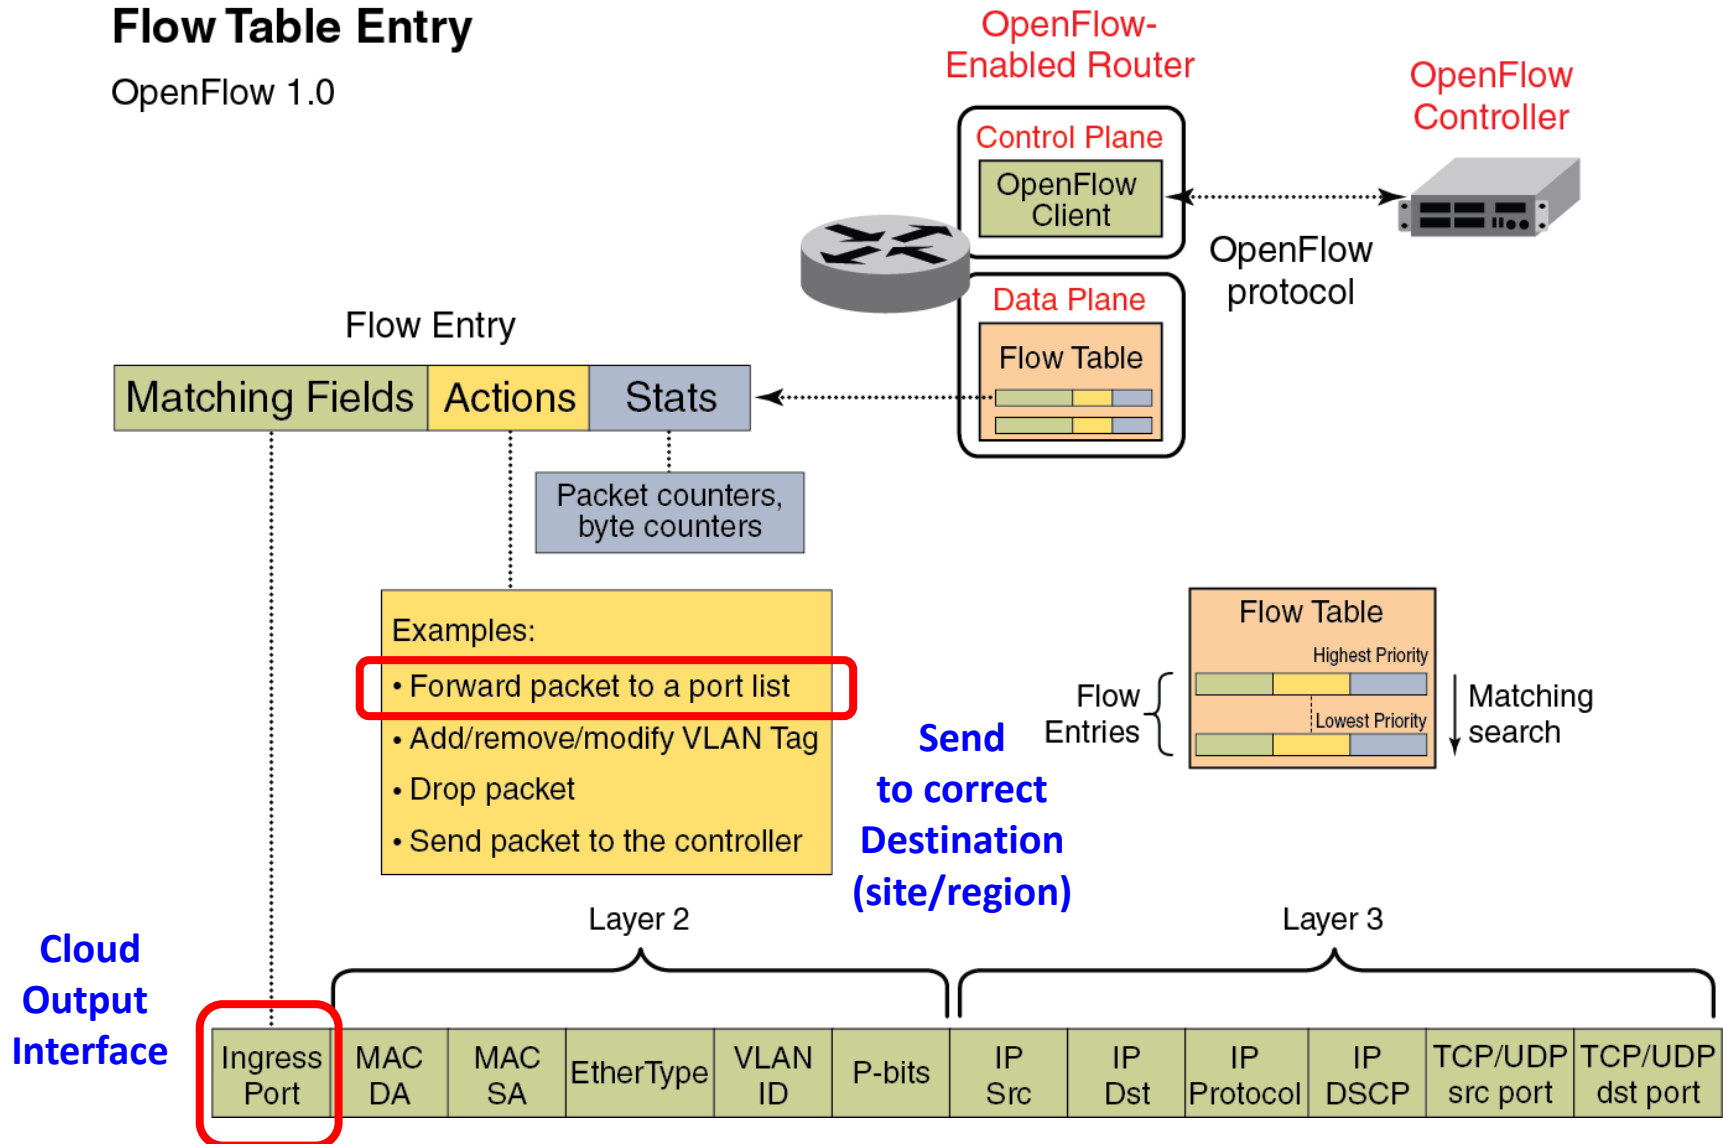
```
fvctl-json --passwd-file=passwd add-slice OPENSTACK-VLAN-101 tcp:103.22.221.150:6633 user@gmail.com
```

```
fvctl-json --passwd-file=passwd add-flowpace OPENSTACK-VLAN-101-FLOWSPACE 1111111111111181 10 dl_vlan=101 OPENSTACK-VLAN-101=4
```

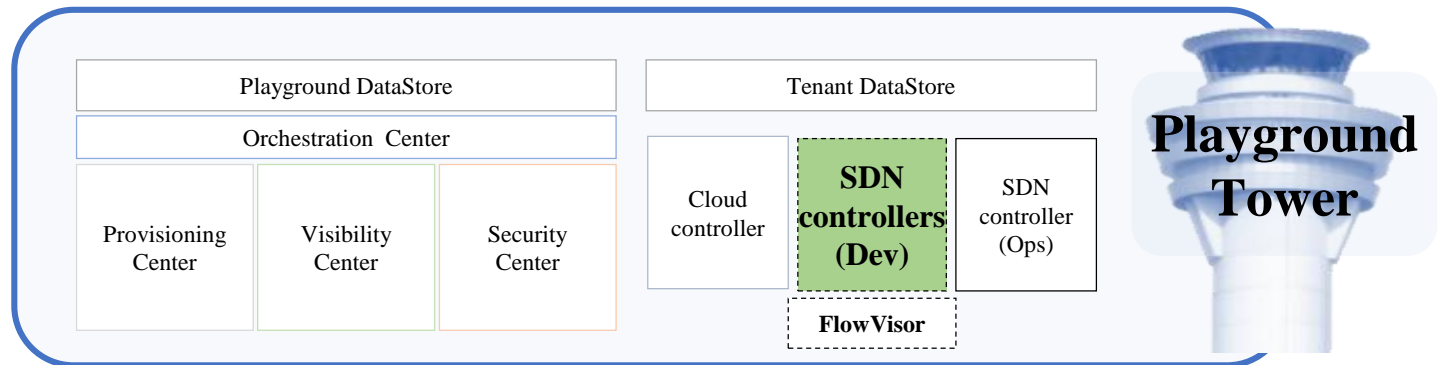
Centralized Control for Developer using Port-based Flow Control (OpenDaylight)

Flow Table Entry

OpenFlow 1.0



Adding Flow Entries using OpenDaylight REST API



GIST Box

```
curl -u admin:admin -H 'Content-type: application/json' -X PUT -d '{"installInHw":"true", "name":"GIST_OUT_ALL", "node": {"id":"11:11:11:11:11:11:11:11", "type":"OF"}, "ingressPort": '1', "priority":"65535", "actions":["OUTPUT=2,3"]}'
```

[http://\\$1:8080/controller/nb/v2/flowprogrammer/default/node/OF/11:11:11:11:11:11:11:11/staticFlow/GIST_OUT_ALL](http://$1:8080/controller/nb/v2/flowprogrammer/default/node/OF/11:11:11:11:11:11:11:11/staticFlow/GIST_OUT_ALL)

```
curl -u admin:admin -H 'Content-type: application/json' -X PUT -d '{"installInHw":"true", "name":"GIST_MY_UM", "node": {"id":"11:11:11:11:11:11:11:11", "type":"OF"}, "ingressPort": '2', "priority":"65535", "actions":["OUTPUT='1'"]}'
```

[http://\\$1:8080/controller/nb/v2/flowprogrammer/default/node/OF/11:11:11:11:11:11:11:11/staticFlow/GIST_MY_UM](http://$1:8080/controller/nb/v2/flowprogrammer/default/node/OF/11:11:11:11:11:11:11:11/staticFlow/GIST_MY_UM)

```
curl -u admin:admin -H 'Content-type: application/json' -X PUT -d '{"installInHw":"true", "name":"GIST_TH_CHULA", "node": {"id":"11:11:11:11:11:11:11:11", "type":"OF"}, "ingressPort": '3', "priority":"65535", "actions":["OUTPUT='1'"]}'
```

[http://\\$1:8080/controller/nb/v2/flowprogrammer/default/node/OF/11:11:11:11:11:11:11:11/staticFlow/GIST_TH_CHULA](http://$1:8080/controller/nb/v2/flowprogrammer/default/node/OF/11:11:11:11:11:11:11:11/staticFlow/GIST_TH_CHULA)

TH-CHULA Box


```
curl -u admin:admin -H 'Content-type: application/json' -X PUT -d '{"installInHw":"true", "name":"TH_OUT_ALL", "node": {"id":"11:11:11:11:11:11:11:81", "type":"OF"}, "ingressPort": '1', "priority":"65535", "actions":["OUTPUT=2"]}'
```

[http://\\$1:8080/controller/nb/v2/flowprogrammer/default/node/OF/11:11:11:11:11:11:11:81/staticFlow/TH_OUT_ALL](http://$1:8080/controller/nb/v2/flowprogrammer/default/node/OF/11:11:11:11:11:11:11:81/staticFlow/TH_OUT_ALL)

```
curl -u admin:admin -H 'Content-type: application/json' -X PUT -d '{"installInHw":"true", "name":"TH_GIST", "node": {"id":"11:11:11:11:11:11:11:81", "type":"OF"}, "ingressPort": '2', "priority":"65535", "actions":["OUTPUT='1'"]}'
```

[http://\\$1:8080/controller/nb/v2/flowprogrammer/default/node/OF/11:11:11:11:11:11:11:81/staticFlow/TH_GIST](http://$1:8080/controller/nb/v2/flowprogrammer/default/node/OF/11:11:11:11:11:11:11:81/staticFlow/TH_GIST)

Verify OpenStack Installation




openstack.

Log in

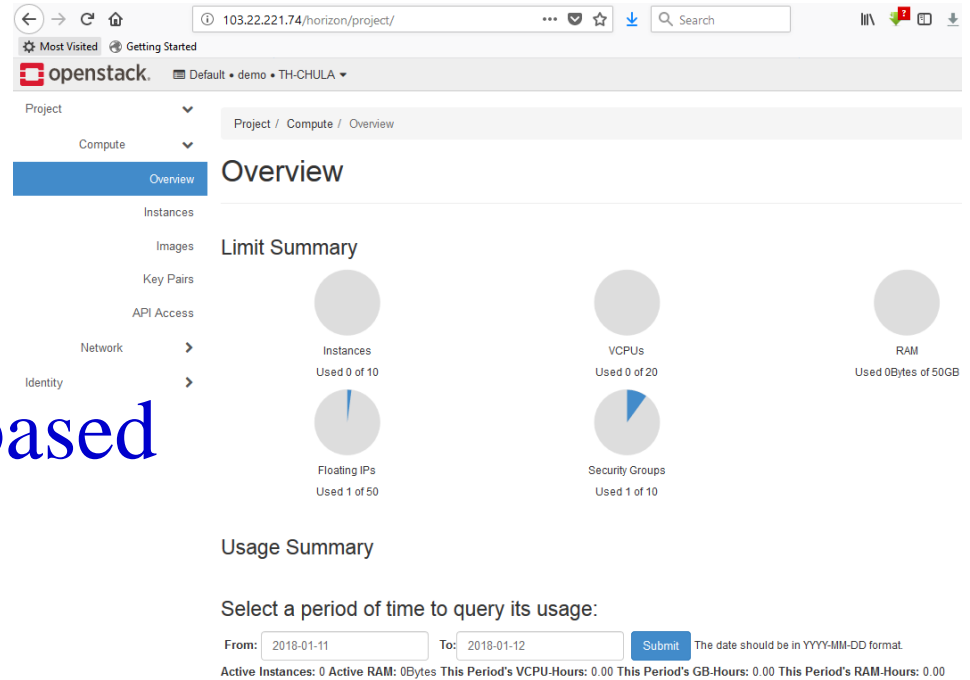
Domain

User Name

Password

Enter OpenStack
Account Info.

Web-based



[SSH Remote Access] Enter System Account

\$ ssh <system account>@103.22.221.74

[After Login]

\$ cd openstack

\$. demo-openrc

\$ openstack --os-region-name TH-CHULA server list

API-based

```
Using username "tein".
Welcome to Ubuntu 16.04.3 LTS (GNU/Linux 4.10.0-42-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

14 packages can be updated.
8 updates are security updates.

*** System restart required ***
Welcome to the perfSONAR toolkit v4.0.2.1-2

You may create accounts to manage this host through the web interface by running the following as root:

/usr/lib/perfsonar/scripts/mptoolkit-configure.py

The web interface should be available at:

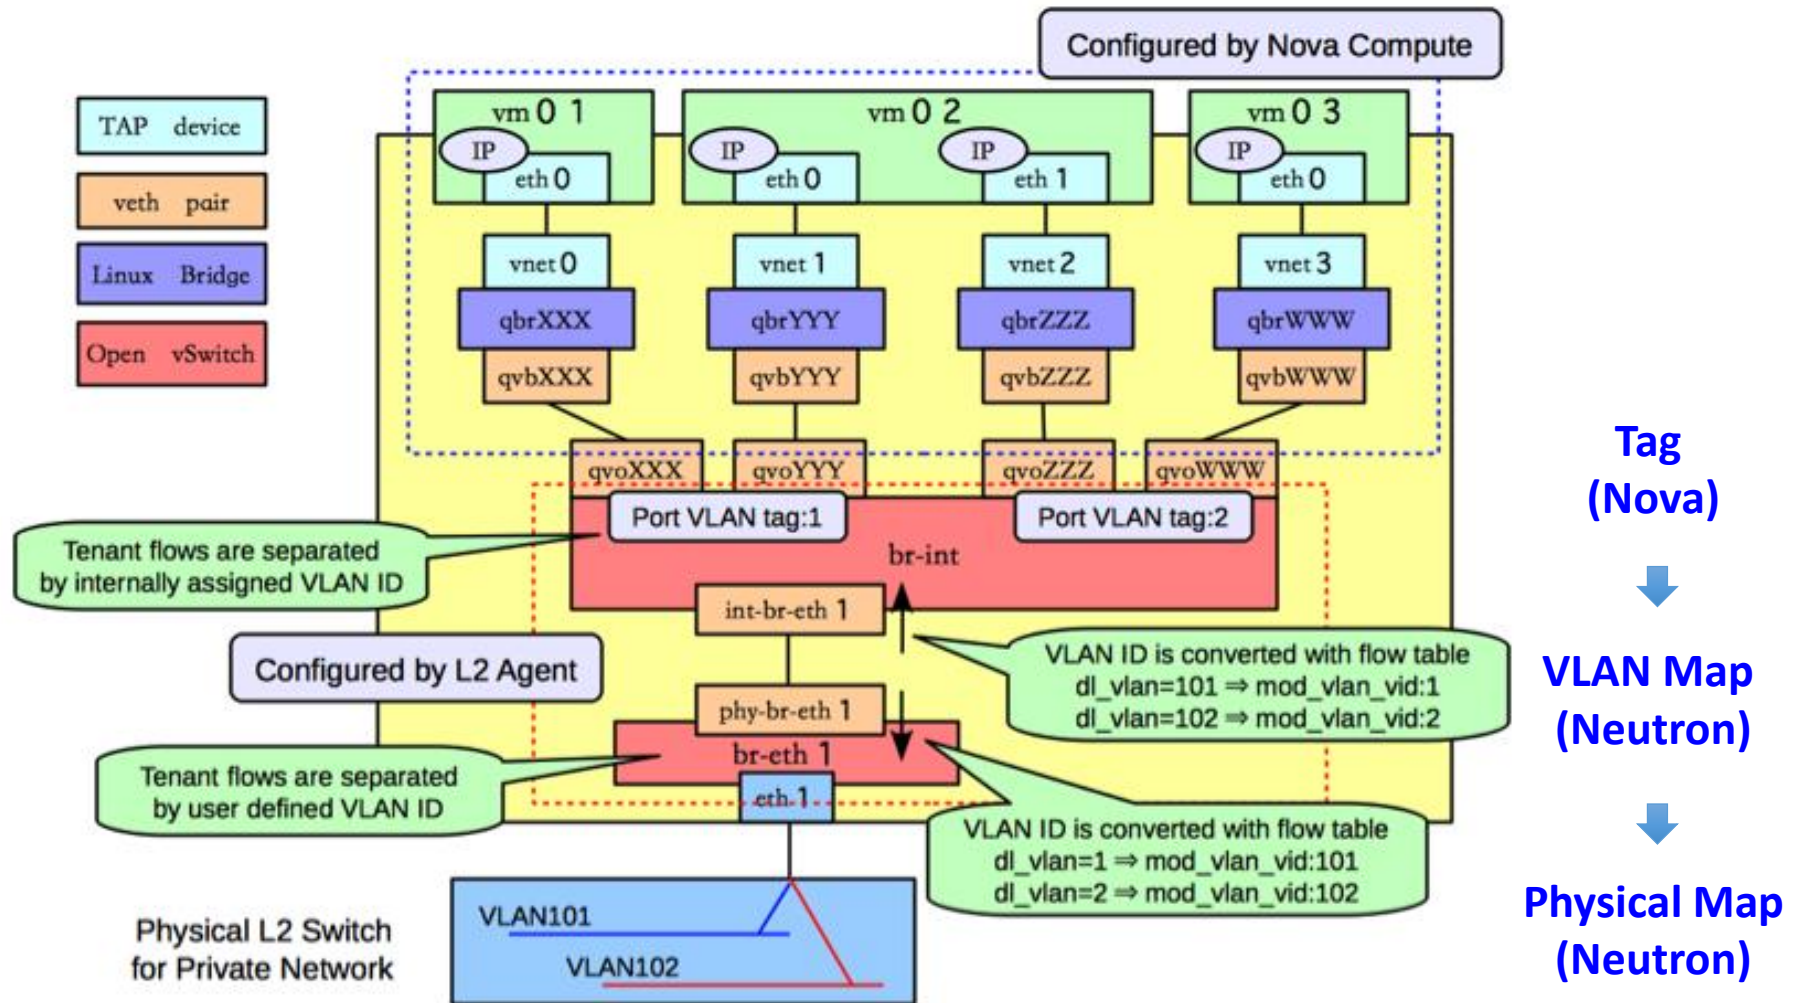
https://161.200.25.99/toolkit
Last login: Fri Jan 12 00:42:07 2018 from 209.237.53.67
tein@Server-Box-CHULA:~$ ssh netcs@103.22.221.74
Welcome to Ubuntu 16.04.3 LTS (GNU/Linux 4.10.0-38-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

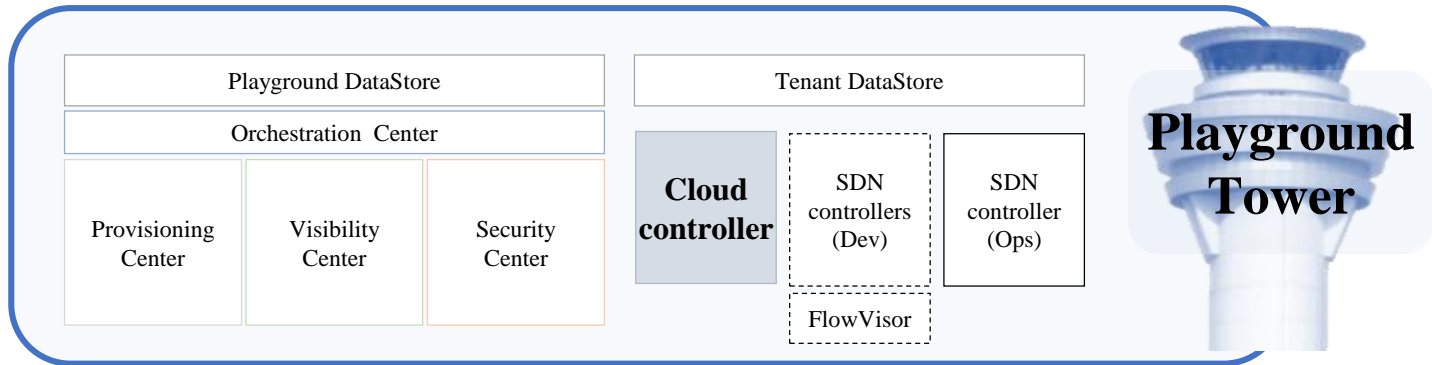
21 packages can be updated.
17 updates are security updates.

*** System restart required ***
Last login: Fri Jan 12 14:42:14 2018 from 161.200.25.99
netcs@S-Center:~$ cd openstack/
netcs@S-Center:~/openstack$ . demo-openrc
netcs@S-Center:~/openstack$
netcs@S-Center:~/openstack$
netcs@S-Center:~/openstack$ openstack --os-region-name TH-CHULA server list
+-----+-----+-----+-----+-----+
| ID                                           | Name                               | Status | Networks | Image Name |
+-----+-----+-----+-----+-----+
| 3d0d73f6-30f9-4363-aefa-d70ff19565e5      | chula-ops-vm                      | ACTIVE | datapath101=192.168.1.138; control101=10.1.101.12 | Xenial-Ubuntu 16.04 |
```

SDN-Cloud Tenant Networking VLAN-based Tenant Network (OpenStack Neutron)



Creating OpenStack Neutron Networks



```
USER="demo"
PROJECT="demo"
REGION1="TH-CHULA"
NET_ID=01
```

```
. admin-openrc
```

```
PROJECT_ID=`openstack project list | grep $PROJECT | awk '{print $2}`
USER_ID=`openstack user list | grep $USER | awk '{print $2}`
```

Create Data Network

```
neutron net-create --os-region-name $REGION1 --tenant-id $PROJECT_ID --provider:network_type vlan --provider:physical_network provider --provider:segmentation_id 1$NET_ID datapath1$NET_ID
```

Create Control Network

```
neutron net-create --os-region-name $REGION1 --tenant-id $PROJECT_ID --provider:network_type vxlan control1$NET_ID
```

Create Subnets

```
. $USER-openrc
openstack --os-region-name $REGION1 subnet create --os-project-id $PROJECT_ID datapath1$NET_ID --network datapath1$NET_ID --subnet-range 192.168.1.0/24 --allocation-pool start=192.168.1.1,end=192.168.1.15 --no-dhcp --gateway none
openstack --os-region-name $REGION1 subnet create --os-project-id $PROJECT_ID --network control1$NET_ID control1$NET_ID --subnet-range 10.1.1$NET_ID.0/24 --dhcp --gateway 10.1.1$NET_ID.1 --dns-nameserver 8.8.8.8
```

Create Virtual Router

```
openstack --os-region-name $REGION1 router create --os-project-id $PROJECT_ID router1$NET_ID
openstack --os-region-name $REGION1 router add subnet router1$NET_ID `openstack --os-region-name $REGION1 network list | grep control1$NET_ID | awk '{print $6}`
openstack --os-region-name $REGION1 router set router1$NET_ID --external-gateway public
```

Creating OpenStack Instance

103.22.221.74/horizon/project/instances/?action=rov ... ☆ ⬇ 🔍 Search

Most Visited Getting Started

openstack. Default • demo • TH-CHULA demo

Project / Compute / Instances

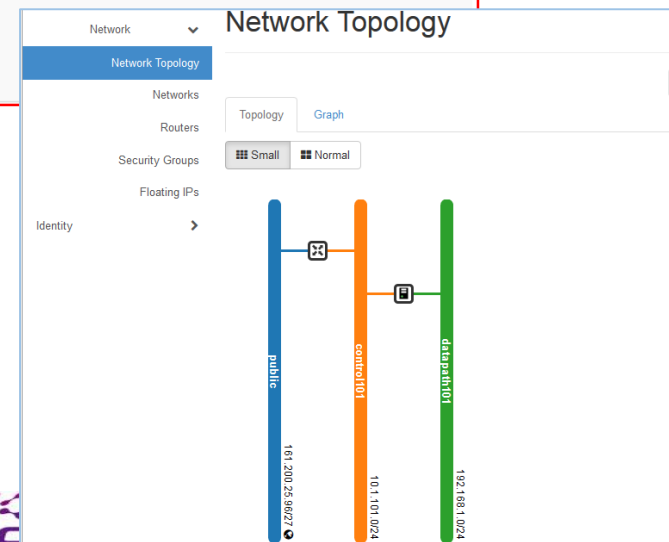
Instances

Instance ID = Filter Launch Instance Delete Instances More Actions

Displaying 1 item

<input type="checkbox"/>	Instance Name	Image Name	IP Address	Flavor	Key Pair	Status	Availability Zone	Task	Power State	Time since created	Actions
<input type="checkbox"/>	chula-ops-vm	Xenial-Ubuntu 16.04	control101 10.1.101.12 Floating IPs: 161.200.25.109	m1.small	chula-keypair	Active	nova	None	Running	2 minutes	Create Snapshot

Displaying 1 item



Access VM, Configure Datapath Interface & Verify end-to-end Connectivity

Directly access VM

```
netcs@S-Center:~/openstack/openstack-keys$ ssh -i chula-keypair.pem ubuntu@161.200.25.109
The authenticity of host '161.200.25.109 (161.200.25.109)' can't be established.
ECDSA key fingerprint is SHA256:VMjgolp495DIDkBS873IhtcthllyU8zntWkclGGPN5o.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '161.200.25.109' (ECDSA) to the list of known hosts.
Welcome to Ubuntu 16.04.3 LTS (GNU/Linux 4.4.0-96-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

Get cloud support with Ubuntu Advantage Cloud Guest:
http://www.ubuntu.com/business/services/cloud

0 packages can be updated.
0 updates are security updates.
```

Show network interfaces

```
ubuntu@chula-ops-vm:~$ ip link show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: ens3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc pfifo_fast state UP mode DEFAULT group default qlen 1000
    link/ether fa:16:3e:62:43:d0 brd ff:ff:ff:ff:ff:ff
3: ens4: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN mode DEFAULT group default qlen 1000
    link/ether fa:16:3e:c6:2d:62 brd ff:ff:ff:ff:ff:ff
```

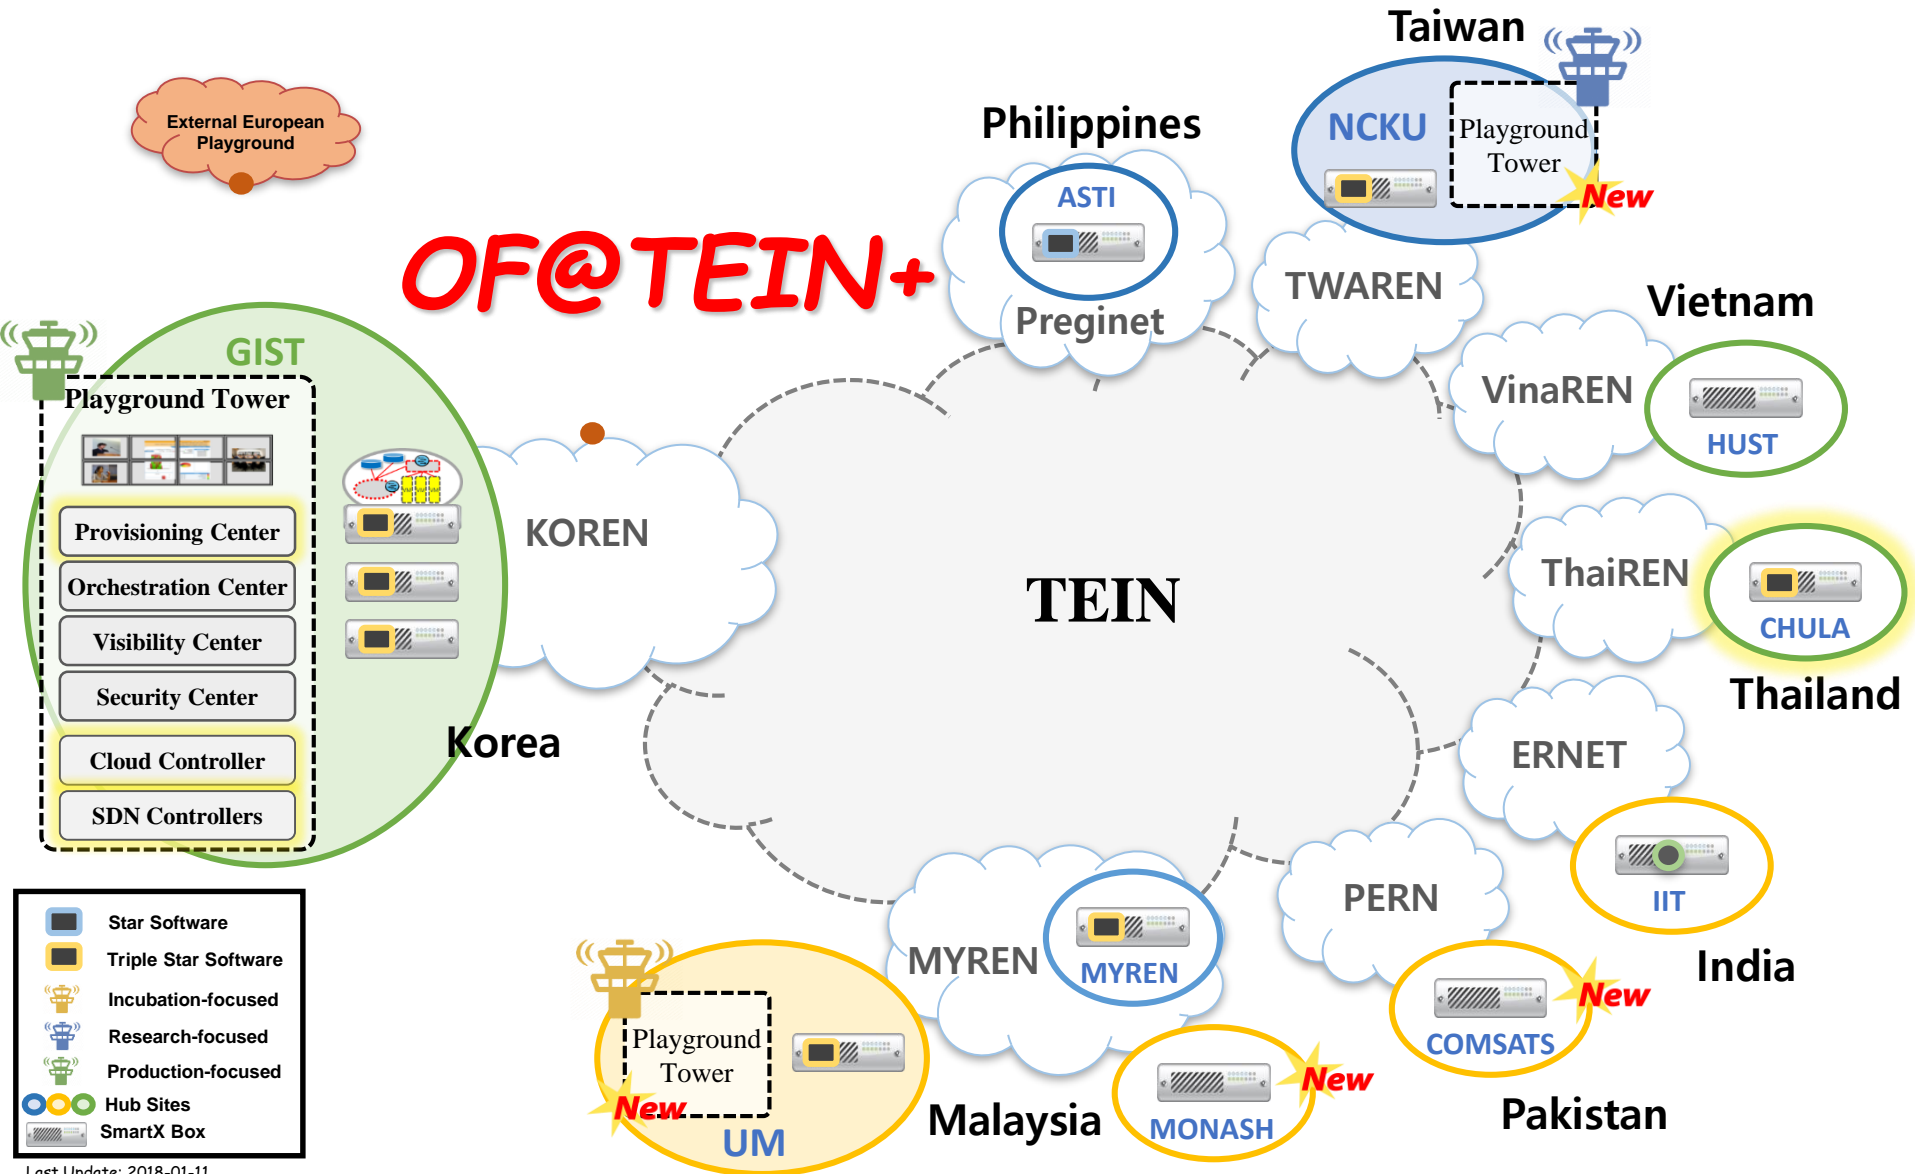
Configure datapath interface

```
ubuntu@chula-ops-vm:~$ sudo ifconfig ens4 192.168.1.138 netmask 255.255.255.0 up
sudo: unable to resolve host chula-ops-vm
ubuntu@chula-ops-vm:~$
ubuntu@chula-ops-vm:~$
ubuntu@chula-ops-vm:~$
ubuntu@chula-ops-vm:~$
ubuntu@chula-ops-vm:~$ ping 192.168.1.7
PING 192.168.1.7 (192.168.1.7) 56(84) bytes of data.
64 bytes from 192.168.1.7: icmp_seq=1 ttl=64 time=358 ms
64 bytes from 192.168.1.7: icmp_seq=2 ttl=64 time=97.4 ms
```

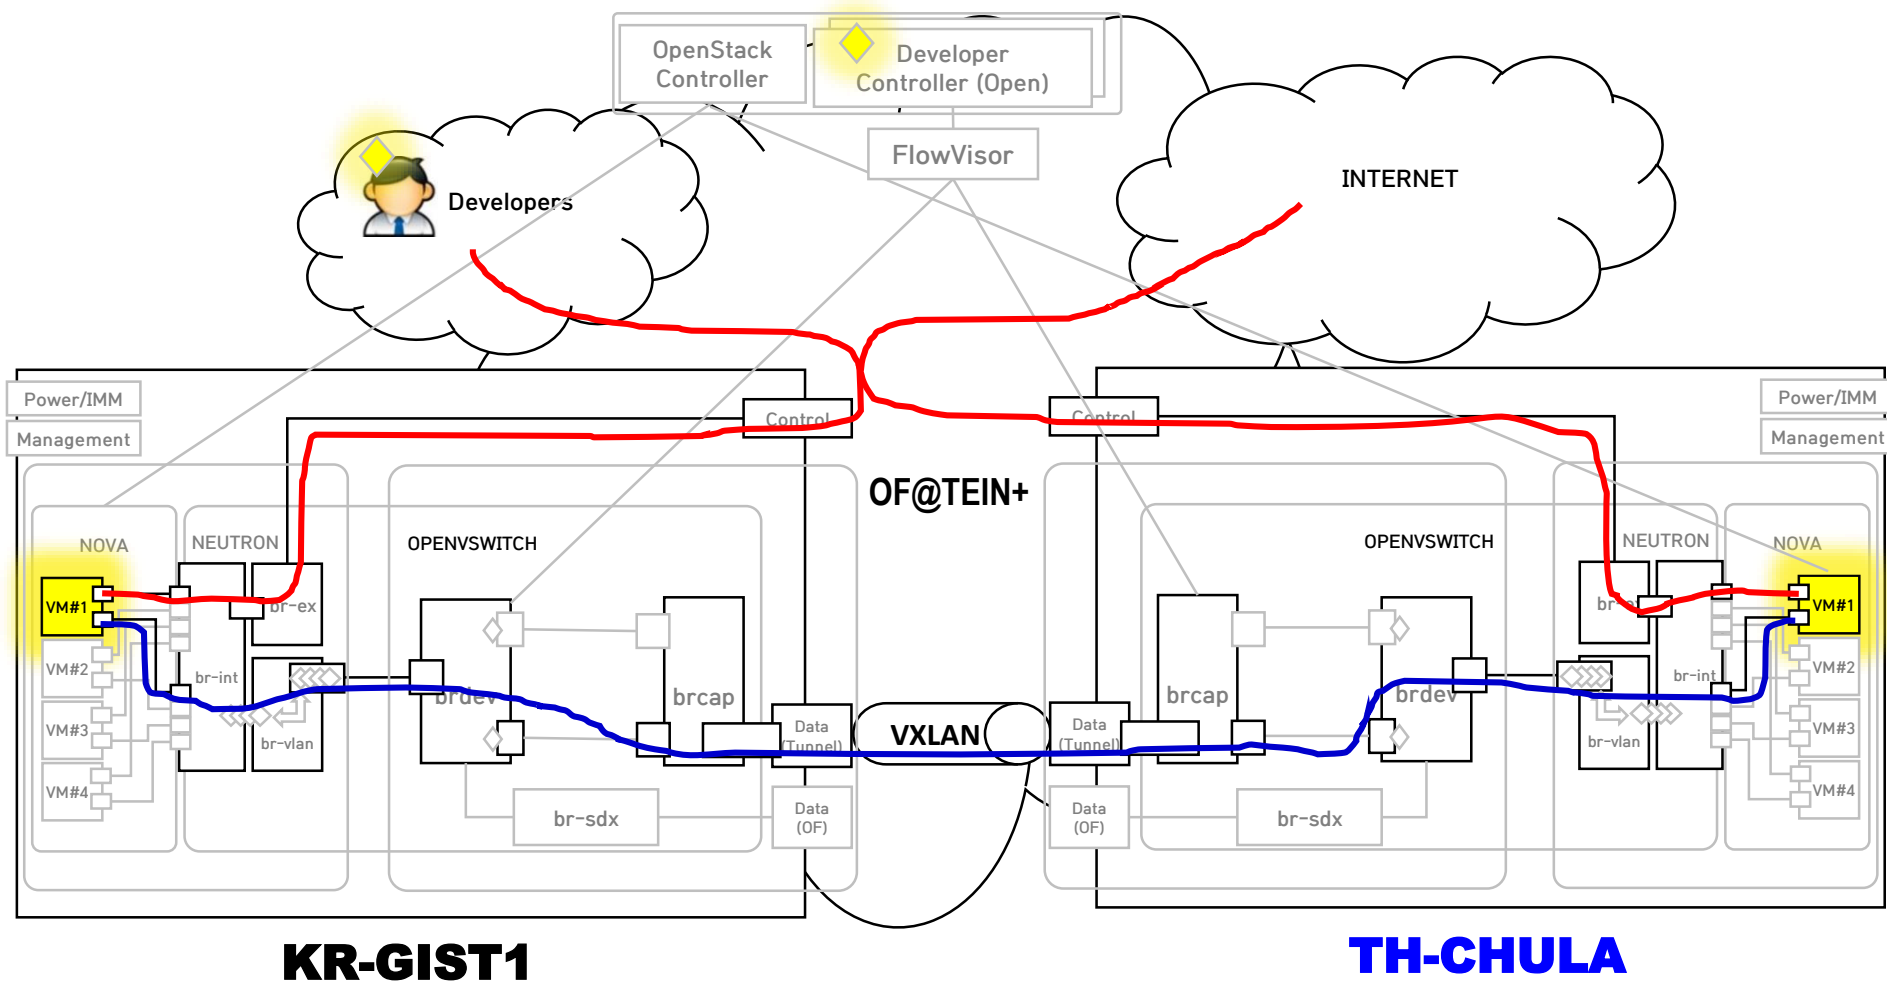
Access another VM in same tenant

SmartX Box & Inter-connect: Provisioning Demonstration

OF@TEIN+ Playground: **Upgrading** (Existing) / Adding (New) Site into Multi-site Cloud



SmartX Box & Inter-connect: Provisioning Verification Scenario



Thank You

ops@oftein.net